

Hybrid Retrieval for Document QA

We evaluate hybrid retrieval combining sparse and dense signals. Experiments use a fixed corpus of 12k pages with automatic layout parsing.
Mean latency was 412 ms per query on CPU-only hardware.
Primary metric: $F1 = 2 * (precision * recall) / (precision + recall)$.

Results

Hybrid retrieval achieved 78.3% answer accuracy on the development split.




Figure 1: Ablation over retriever components. Hybrid fusion improves recall@8.

Table 1 — Benchmark rows

System	F1	Latency (ms)
Hybrid	0.78	412